MultiIFEval: An Instruction Following Benchmark in 304 Languages

Anonymous ACL submission

Abstract

We introduce MultiIFEval, a large-scale instruction-following benchmark that spans 304 languages. Starting from the English IFEval dataset, we translate and localise each prompt–response pair into 304 target languages using Gemini-3-flash-preview.

We also create a new manually translated and localised Danish variant of the dataset, and conduct manual expert validation on the quality difference between the machine translated version of the dataset and the manually translated version, which showed that TODO. We also evaluated models on both the machine translated and manually translated versions of the Danish, Estonian and French versions of the dataset, showing that TODO.

Lastly, we evaluate XXX models on all 304 languages of the dataset, showing that TODO.

1 Introduction

Instruction following evaluation has become central to developing practical AI systems capable of understanding and executing natural language commands. Early benchmarks such as IFEval (Zhou et al., 2023) established a standardised framework for measuring how well models can follow instructions in English, but they quickly revealed a critical limitation: the lack of multilingual coverage. M-IFEval (Li et al., 2025), XIFBench (He et al., 2024), and Multi-IF (He et al., 2024) have since expanded the scope to include 3, 6 and 8 languages, respectively. Marco-Bench-MIF (Zeng et al., 2025) extended the work to a commendable 30 languages, but these efforts are still limited to a small subset of the world's languages.

To bridge this gap, we introduce **MultiIFEval**, a benchmark that expands IFEval to 304 target languages. Leveraging a large language model, we both translate and localise every English prompt–response pair into each language, with localised grounding coming from Wikipedia articles in the given language. The resulting dataset spans high-resource languages like Spanish and Mandarin as well as low-resource ones such as Xhosa and Quechua, providing an unprecedented testbed for multilingual instruction following.

In addition to automated translations, we provide a manually translated Danish subset curated by native speakers and validated by experts. This dual approach allows us to quantify the impact of translation quality on model performance—an aspect largely unexplored in prior work. Similar manually translations and localisations have been done for Estonian (?) and French (?) as well, which we also include in our analysis.

Our main contributions are as follows:

- We create a multilingual instruction-following benchmark spanning 304 languages.
- We provide a manually translated and localised Danish subset for expert validation and comparison with machine translations.
- We analyse the model performance discrepancies of the machine translated and manually translated versions of the Danish, Estonian and French versions of the dataset.
- We evaluate XXX models on all 304 languages of the dataset.

2 Dataset Generation

Missing

2.1 Manual Danish Translation

Missing

2.2 Machine Translated Dataset

Missing

3 Evaluations of Language Models

Missing

4 Comparison of Machine Translated and Manually Translated Datasets

Missing

5 Conclusion

Missing

6 Acknowledgements

Missing

References

Yun He, Di Jin, Chaoqi Wang, Chloe Bi, Karishma Mandyam, Hejia Zhang, Chen Zhu, Ning Li, Tengyu Xu, Hongjiang Lv, and 1 others. 2024. Multi-if: Benchmarking llms on multi-turn and multilingual instructions following. *arXiv preprint arXiv:2410.15553*.

Zhenyu Li, Kehai Chen, YUNFEI LONG, Xuefeng Bai, Yaoyin Zhang, Xuchen Wei, Juntao Li, and Min Zhang. 2025. Xifbench: Evaluating large language models on multilingual instruction following. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*.

Bo Zeng, Chenyang Lyu, Sinuo Liu, Mingyan Zeng, Minghao Wu, Xuanfan Ni, Tianqi Shi, Yu Zhao, Yefeng Liu, Chenyu Zhu, Ruizhe Li, Jiahui Geng, Qing Li, Yu Tong, Longyue Wang, Weihua Luo, and Kaifu Zhang. 2025. Marco-bench-MIF: On multilingual instruction-following capability of large language. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 24058–24072, Vienna, Austria. Association for Computational Linguistics.

Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. 2023. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*.

A Example Appendix

Missing